

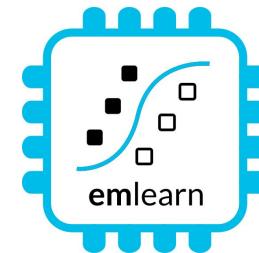
# Microcontrollers and Machine Learning

with MicroPython and emlearn

<https://github.com/emlearn/emlearn-micropython>



Jon Nordby [jononor@gmail.com](mailto:jononor@gmail.com)  
foss-north 2025





We utilise sound and vibration analysis to detect and warn you of upcoming errors in your technical infrastructure before they happen.

 **sound sensing**

## Condition Monitoring

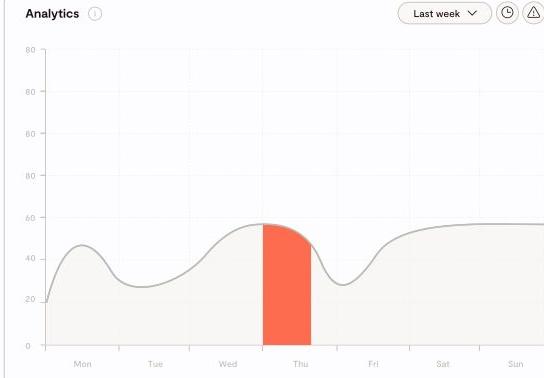
Devices overview Search Filter by Organization Hotel Manana

**Tech\_Storage**

Location	Room	Status
Sandstuvein...	Island 04	<span>Green</span>
Storgata 25...	TBJ 291	<span>Yellow</span>
Chr. Krohgs...	Tech 275	<span>Red</span>
Møllergata 12...	Ohio 3	<span>Green</span>
Ruseløkkveien...	HKM 261	<span>Green</span>
Kristian IVs...	Freeway 273	<span>Yellow</span>
C. J. Hambro...	Oxaca2	<span>Green</span>
Akersgata 65...	Storage_3	<span>Green</span>

**Tech 275**

**Analytics** Last week More Less



Mon Tue Wed Thu Fri Sat Sun

Trusted by Nordic market leaders

# Goal

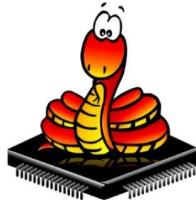
Purpose of this presentation

You, as a **Python developer**,  
can **build an IoT sensor**  
on a **microcontroller**  
using **MicroPython**

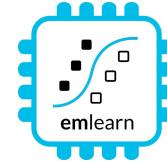
+ Including  
**on-edge Machine Learning**  
with **emlearn-micropython**

# Python on Microcontroller

Jumping right into it



**MicroPython**



# What is a microcontroller?

Modern microcontroller:  
A complete programmable System-on-Chip

Example: ESP32-S3FH4R2

**32 bit CPU, 240 Mhz**

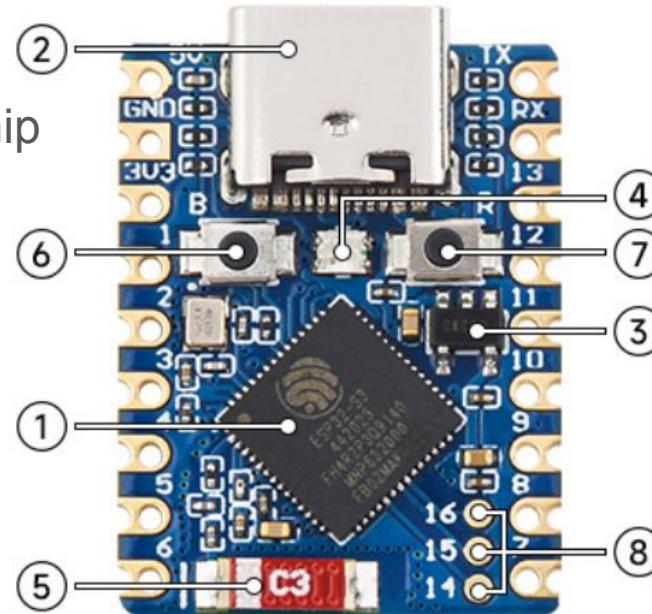
**Floating Point Unit**

**2 MB RAM**

**4 MB FLASH**

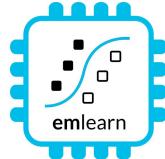
**WiFi**

**Bluetooth Low Energy**  
**USB-C**



Espressif ESP32-S3FH4R2 chip:  
Waveshare ESP32-S3-Tiny board:

2.5	USD
6	USD

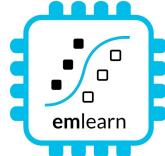


# Hardware tinkering optional - complete devices available



Complete ESP32 based device with sensors etc.:

20 - 50 USD



# Installing MicroPython

Download prebuilt firmware

<https://micropython.org/download/?port=esp32>

Flash firmware to device

pip install esptool

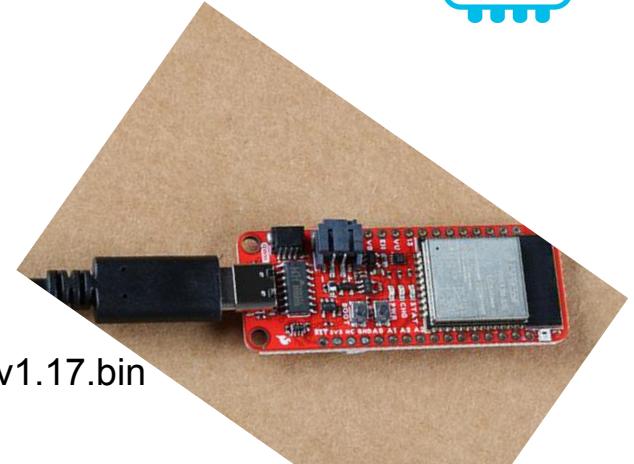
esptool.py --chip esp32 --port ... erase\_flash

esptool.py --chip esp32 --port ... write\_flash -z 0 micropython-v1.17.bin

Connect to device

pip install mpremote

mpremote repl



```
MicroPython v1.8.3-24-g095e43a on 2016-08-16; ESP module
Type "help()" for more information.
>>> print('Hello world!')
Hello world!
>>> █
```

IDE (optional): Viper IDE, Thonny, et.c.

# Temperature sensor - hardware



**Example**  
MPU6886 accelerometer  
built-in temperature sensor

**Temperature Sensor**  
DS1820  
Water-proof

**Room Monitoring**  
+ CO2  
+ humidity  
+ temperature

# Temperature sensor - code

Using <https://viper-ide.org/>

Zero-install. Connect to device via USB

1. Read the sensor in a loop
2. Send data using MQTT
3. Wait until next measurement

The same approach can be used for  
sensing other slow changing phenomena

```
test.py • X
1 from mqtt_as import MQTTClient, config
2 import asyncio
3 from mpu6886 import MPU6886
4 from machine import I2C
5
6 # Local configuration
7 config['ssid'] = 'FIXME' # Optional on ESP8266
8 config['wifi_pw'] = 'FIXME'
9 config['server'] = 'test.mosquitto.org'
10
11 mpu = MPU6886(I2C(0, sda=21, scl=22, freq=100000))
12
13 v async def main(client):
14     print('main-start')
15     await client.connect()
16     print('connected')
17
18 v     while True:
19         t = mpu.temperature
20         print('publish-data', t)
21         await client.publish('pydataglobal2024/send', f'{t:.2f}', qos = 0)
22         await asyncio.sleep(30)
23
24 MQTTClient.DEBUG = True # Optional: print diagnostic messages
25 client = MQTTClient(config)
26 v try:
27     asyncio.run(main(client))
28 v finally:
29     client.close()
```

Using [peterhinch/micropython-mqtt](https://github.com/peterhinch/micropython-mqtt)  
and [jonnor/micropython-mpu6886](https://github.com/jonnor/micropython-mpu6886)

# Installing packages

MicroPython has a package manager “**mip**”

Directly on device!

ViperIDE supports install via UI

Custom packages via URL

The screenshot shows the ViperIDE interface. On the left, there's a sidebar with 'Tools' and 'Python' sections. Under 'Tools', it lists 'device connection' (WebREPL over internet, WebREPL in local network, Bluetooth REPL, P2P Bridge), 'Device Benchmark', and 'package manager'. Under 'package manager', there are 'Install package via link' (selected), 'Prettify current file', and 'Minify current file'. A modal dialog box is open, prompting 'Enter package name or URL:' with the value 'https://github.com/jonnor/micropython-mpu6886/'. Below the modal, the code editor displays a Python script:

```
7 config['ssid'] = 'FIXME' # Optional on ESP8266
8 config['wifi_pw'] = 'FIXME'
9 config['server'] = 'test.mosquitto.org'
10
11 mpu = MPU6886(I2C(0, sda=21, scl=22, freq=100000))
12
13 async def main(client):
14     print('main-start')
15     await client.connect()
16     print('connected')
17
18     while True:
19         t = mpu.temperature
20         print('publish-data', t)
21         await client.publish('pydataglobal2024/send', f'{t:.2f}', qos = 0)
22         await asyncio.sleep(30)
23
24 MQTTClient.DEBUG = True # Optional: print diagnostic messages
25 client = MQTTClient(config)
26 try:
27     asyncio.run(main(client))
28 finally:
29     client.close()
```

At the bottom, a terminal window shows the MicroPython prompt:

```
>>> Type "help()" for more information.
>>> MicroPython v1.23.0-3.g2ceccb531.dirty on 2024-09-25; Generic ESP32 module with SPIRAM with ESP32
Type "help()" for more information.
>>> >ESP32 module with SPIRAM with ESP32
Type "help()" for more information.
>>> summary-queue-overflow
summary-queue-overflow
```

# Sensor nodes with MicroPython

What can be done on an ESP32  
microcontroller with MicroPython  
and how to make it work

# Temperature logger

Temperature  
0.1 Hz

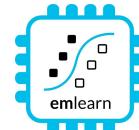


Feasibility

Pure Python

# Activity tracker

Accelerometer  
100 Hz



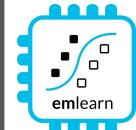
Random Forest  
classifier  
emlearn\_trees

# Noise monitor

Microphone  
16000 Hz

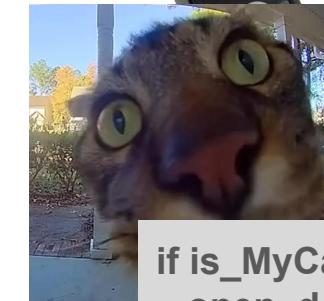


Infinite Impulse  
Response filters  
emlearn\_iir

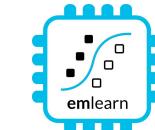


# Image Classifier

Camera  
27000 bytes/s



```
if is_MyCat(img):  
    open_door()
```



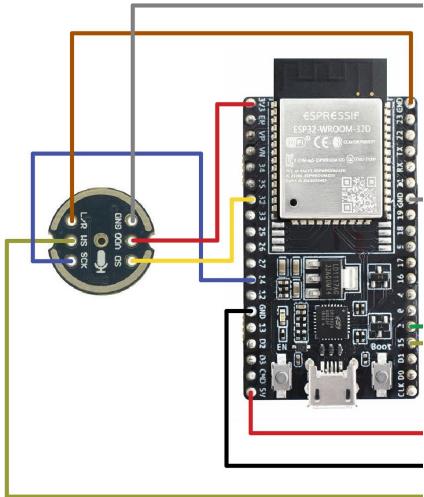
Convolutional  
Neural Network  
emlearn\_cnn

Python with C modules

Sound  
sensor

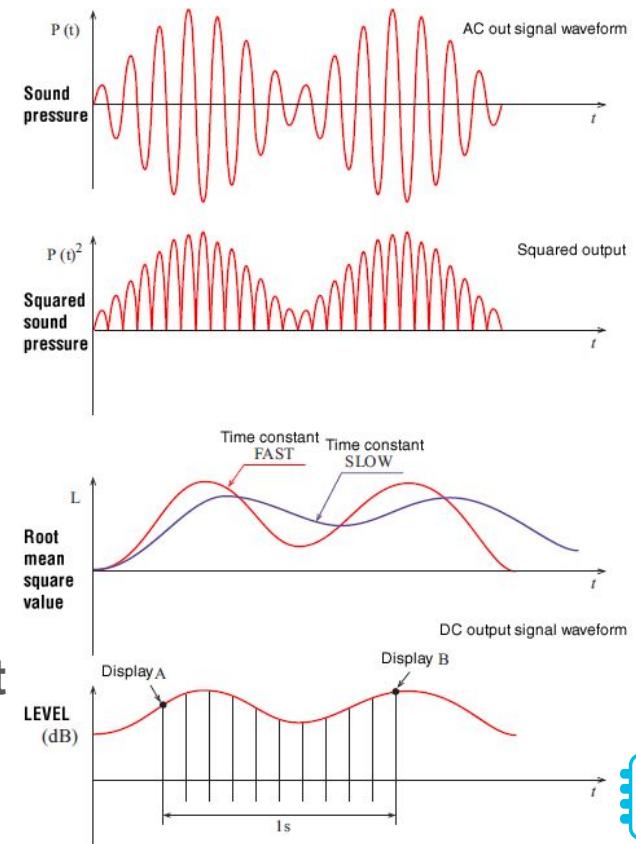
# Sound sensor

## I2S digital microphone (Example wiring)



**Complete hardware unit**  
LilyGo TTGO  
T-Camera S3

## Soundlevel calculation Processing steps



# Sound sensor - Install it

[https://github.com/emlearn/emlearn-micropython/tree/master/examples/soundlevel\\_iir](https://github.com/emlearn/emlearn-micropython/tree/master/examples/soundlevel_iir)

## Running on device (ViperIDE)

The fastest and easiest way to install on your device is to use Viper IDE. This will install the library and the example code automatically.

 Run in ViperIDE

In Viper IDE, you can select which example file to run (described below), and hit the Play button to run it.

```
# I2S audio input
from machine import I2S
audio_in = I2S(0, sck=Pin(26), ws=Pin(32), sd=Pin(33),
               mode=I2S.RX, bits=16, format=I2S.MONO, rate=16000,
               )
# allocate sample arrays
chunk_samples = int(AUDIO_SAMPLERATE * 0.125)
mic_samples = array.array('h', (0 for _ in range(chunk_samples))) # int16
# memoryview used to reduce heap allocation in while loop
mic_samples_mv = memoryview(mic_samples)
# global to share state between callback and main
soundlevel_db = 0.0

meter = SoundlevelMeter(buffer_size=chunk_samples, samplerate=16000)

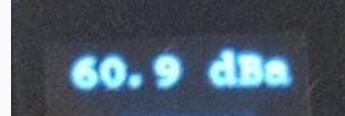
def audio_ready_callback(arg):
    # compute soundlevel
    global soundlevel_db
    soundlevel_db = meter.process(mic_samples)
    # re-trigger audio callback
    _ = audio_in.readinto(mic_samples_mv)

def main():
    # Use Non-Blocking I/O with callback
    audio_in.irq(audio_ready_callback)
    # Trigger first audio readout
    audio_in.readinto(mic_samples_mv)

    while True:
        render_display(db=soundlevel_db)
        time.sleep_ms(200)

if __name__ == '__main__':
    main()
```

Microphone input  
**machine.I2S**



Rendering text/widgets to screen  
<https://github.com/peterhinch/micropython-nano-gui>

# Audio filters using emlearn\_iir

Standard sound level measurements are  
**A-weighted**. To approximate human hearing.

Implemented using  
**Infinite Impulse Response (IIR) filters**.

Must be computed within 125 ms

Python implementation  
**1100 ms 900% CPU**

emlearn\_iir C module  
**30 ms 20% CPU**

```
import emlearn_iir

# setup
coefficients = array.array('f', [
    1.0383002230320646, 0.0, 0.0, 1.0, -0.016647242439959593, 6.9,
    1.0, -2.0, 1.0, 1.0, -1.7070508390293027, 0.7174637059318595,
    1.0, -2.0, 1.0, 1.0, -1.9838868447331497, 0.9839517531763131
])
iir = emlearn_iir.new(coefficients)

# process
iir.run(samples)
```

```
# Use C module for data conversion
from emlearn_arrayutils import linear_map

def int16_to_float(inp, out):
    return linear_map(inp, out, -2**15, 2**15, -1.0, 1.0)

def float_to_int16(inp, out):
    return linear_map(inp, out, -1.0, 1.0, -2**15, 2**15)
```

# C modules

Written in C.

Defines a Python module with API.  
functions/classes etc.

Can be implemented by users,  
libraries or be part of MicroPython  
core.

Can be portable or specific to one  
hardware/platform

```
// Include the header file to get access to the MicroPython API
#include "py/dynruntime.h"

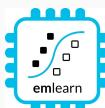
// Helper C function to compute factorial
static mp_int_t factorial_helper(mp_int_t x) {
    if (x == 0) {
        return 1;
    }
    return x * factorial_helper(x - 1);
}

// DEFINE FUNCTION. Callable from Python
static mp_obj_t factorial(mp_obj_t x_obj) {
    mp_int_t x = mp_obj_get_int(x_obj);
    mp_int_t result = factorial_helper(x);
    return mp_obj_new_int(result);
}
static MP_DEFINE_CONST_FUN_OBJ_1(factorial_obj, factorial);

// MODULE ENTRY
mp_obj_t mpy_init(mp_obj_fun_bc_t *self, size_t n_args, size_t n_kw, mp_obj_t *args) {
    // Must be first, it sets up the globals dict and other things
    MP_DYNRUNTIME_INIT_ENTRY

    // Register function in the module's namespace
    mp_store_global(MP_QSTR_factorial, MP_OBJ_FROM_PTR(&factorial_obj));

    // This must be last, it restores the globals dict
    MP_DYNRUNTIME_INIT_EXIT
}
```

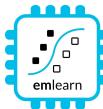


# Native module (.mpy) VS External C module

	<b>Native module</b>	<b>External C module</b>
Installable at runtime	Yes, as .mpy file	No. Must be included in firmware image
Requires SDK/toolchain	No (only to build)	Yes
Code executes from	RAM	FLASH
Limitations	No libc / libm linked * No static BSS *	None
Maturity	Low *	Excellent
Documentation	<a href="https://docs.micropython.org/en/latest/develop/natmod.html">https://docs.micropython.org/en/latest/develop/natmod.html</a>	<a href="https://docs.micropython.org/en/latest/develop/cmodules.html">https://docs.micropython.org/en/latest/develop/cmodules.html</a>

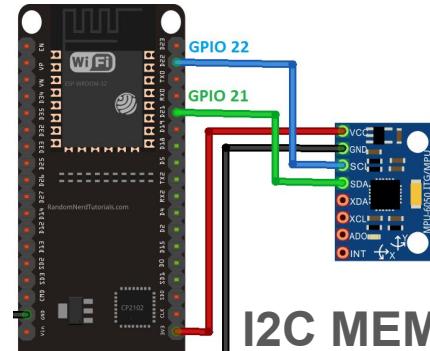
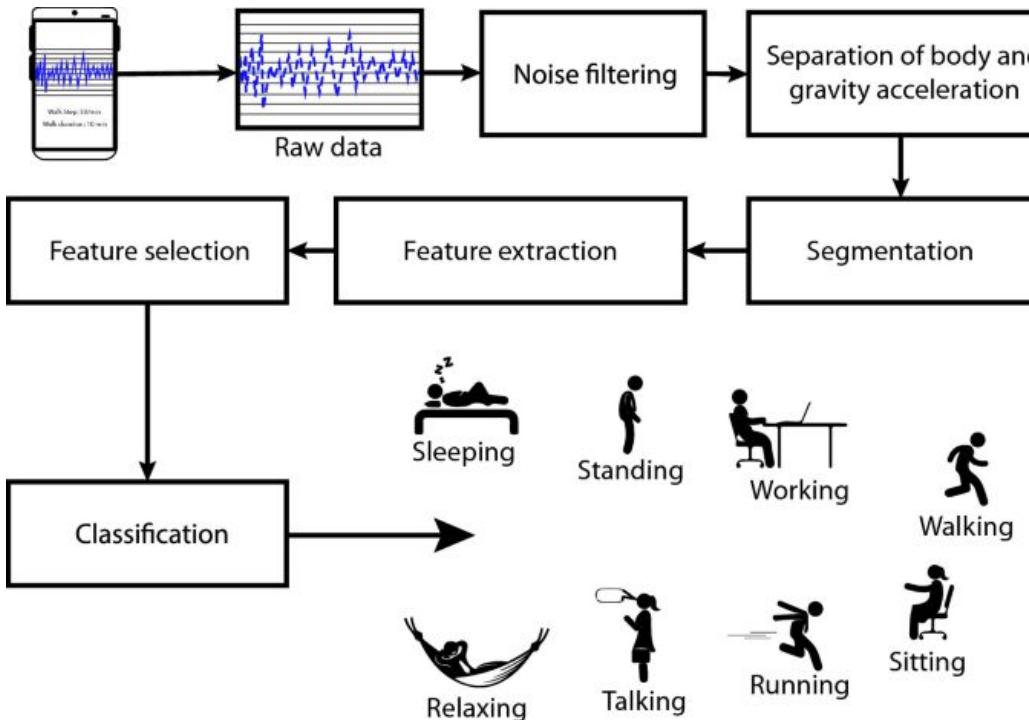
\* Improved greatly in upcoming MicroPython (1.25+).

Contributions by Volodymyr Shymanskyy, Alessandro Gatti, Damien George, and others



# Activity tracker

# Activity tracker - concept



I2C MEMS IMU  
accelerometer/gyro  
(Example wiring)



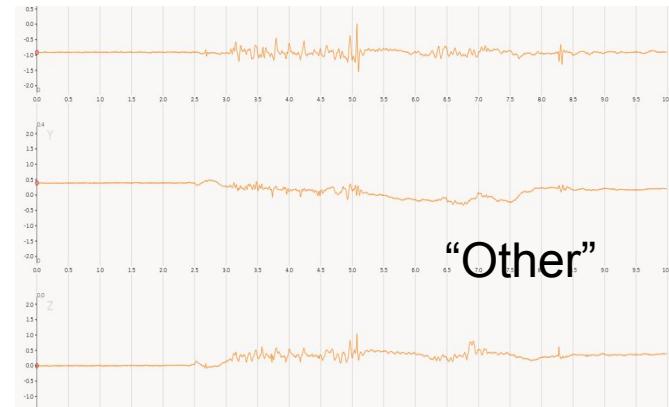
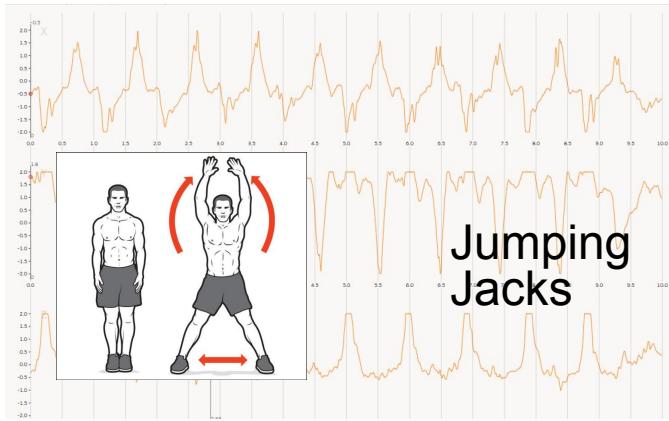
Complete  
hardware unit  
LilyGo  
T-Watch S3

# Recording a dataset



har\_record.py

[https://github.com/emlearn/emlearn-micropython/tree/master/examples/har\\_trees](https://github.com/emlearn/emlearn-micropython/tree/master/examples/har_trees)



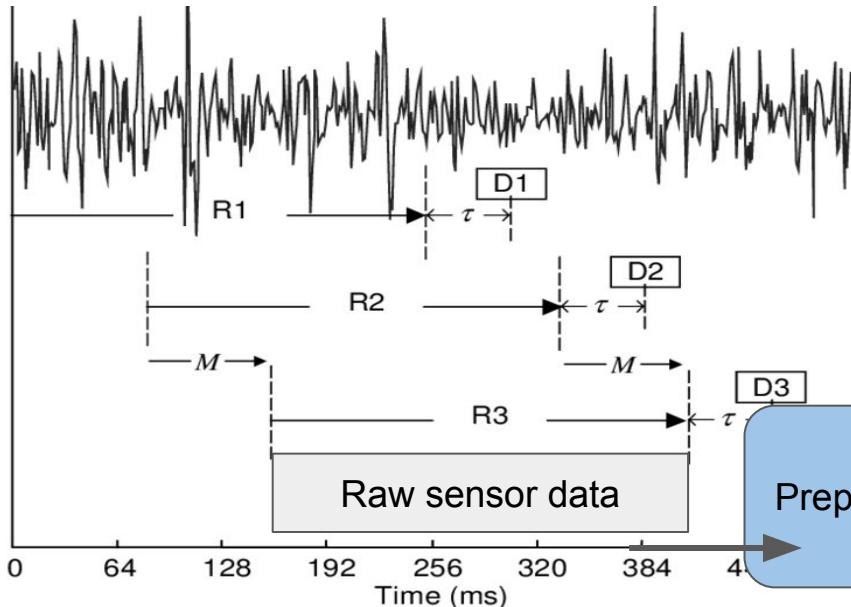
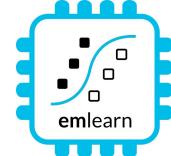
3 separate series,  
1 minute per activity

Saves 10 sec .npy files

Transfer over USB

Cleaned labels in  
Label Studio

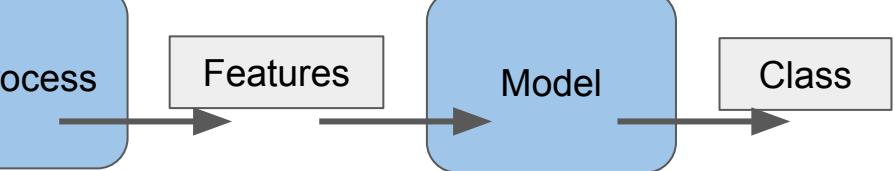
# ML on streams: Continuous classification



The sensor data stream  
is sliced into overlapping windows.  
Each window processed independently

Exercise activity detection:

- 4 second window, every 1 second
- 100 Hz samplerate
- Processing time 200 ms



[ 42, 4002, ..., 329 ]

“Jumping Jacks”

Implementing an IMU/accelerometer/gyro driver? Use the FIFO!  
<https://github.com/orgs/micropython/discussions/15512>

# Activity Tracker - Feature Extraction

Statistical summarizations are useful time-series features, sufficient for basic Human Activity Recognition.

! Preprocessing *must be compatible* between *training* on host PC (CPython) *and device* (MicroPython)

Solution: Write preprocessor for MicroPython, re-use in Python

```
subprocess('micropython preprocess.py data.npy features.npy')
```

Alternative: (when using common MicroPython/CPython subset)

```
import mypreprocessor.py
```

Using micropython-npyfile to read/write Numpy .npy files  
<https://github.com/jonnor/micropython-npyfile/>

```
l = sorted(list(v))
l2 = [x*x for x in l]
sm = sum(l)
sqsum = sum(l2)
avg = sum(l) / len(l)

median = l[MEDIAN]
q25 = l[Q1]
q75 = l[Q3]
iqr = (l[Q3] - l[Q1])

energy = ((sqsum / len(l2)) ** 0.5)
std = ((sqsum - avg * avg) ** 0.5)
```

[https://github.com/emlearn/emlearn-micropython  
/blob/master/examples/har\\_trees/timebased.py](https://github.com/emlearn/emlearn-micropython/blob/master/examples/har_trees/timebased.py)

**Time-based features extraction**  
Are Microcontrollers Ready for Deep Learning-Based Human Activity Recognition?  
Atis Elsts, and Ryan McConville  
<https://www.mdpi.com/2079-9292/10/21/2640>

# Training model on dataset

Using a scikit-learn based pipeline.

```
# Setup subject-based cross validation
splitter = GroupShuffleSplit(n_splits=n_splits, test_size=0.25,
                             random_state=random_state)

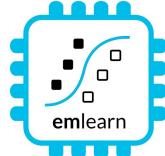
# Random Forest classifier
clf = RandomForestClassifier(random_state = random_state,
                            n_jobs=1, class_weight = "balanced")

# Hyper-parameter search
search = GridSearchCV(clf, param_grid=hyperparameters,
                      scoring=metric, refit=metric, cv=splitter)
search.fit(X, Y, groups=groups)
```

```
(venv) [jon@jon-thinkpad har_trees]$ MIN_SAMPLES_LEAF=150,200,400 python har_train.py
-dataset har_exercise_1 --window-length 400 --window-hop 10
2024-12-04 12:54:52 [info] data-loaded dataset=har_exercise_1
duration=0.016095876693725586 samples=32000
2024-12-04 12:54:56 [info] feature-extraction-done dataset=har_exercise_1
duration=4.534412145614624 labeled_instances=1952 total_instances=1952
Class distribution
activity
jumpingjack    549
lunge           488
other            488
squat            427
Name: count, dtype: int64
Model written to ./har_exercise_1_trees.csv
Testdata written to ./har_exercise_1.testdata.npz
Results
   n_estimators  min_samples_leaf  mean_train_f1_micro  mean_test_f1_micro
0             10              150          0.996311        0.962705
1             10              200          0.995628        0.956557
2             10              400          0.986202        0.920902
```

har\_train.py

```
import emlearn
converted = emlearn.convert(clf)
converted.save(name='gesture', format='csv', file='model.csv')
```



# Deploying trained model

Copy model to device

```
mpremote fs cp gesture_model.csv :
```

Load and run on device

```
import emltrees
model = emltrees.new(10, 1000, 10)
with open('eml_digits.csv', 'r') as f:
    emltrees.load_model(model, f)

features = array.array('h', ...)
out = model.predict(features)
```

## Performance comparison

10 trees, max 100 leaf nodes, “digits” dataset

	Inference time	Program space
m2cgen	60.1 ms	179 kB
everywhereml	17.7 ms	154 kB
<b>emlearn</b>	<b>1.3 ms</b>	<b>15 kB</b>

emlearn is **10x faster and 10x more space efficient** compared to generating Python code



# emlearn for C

## Embedded Friendly

- Portable C99 code
- No dynamic allocations
- Header-only
- High test coverage
- Integer/fixed-point math \*
- Small. 2 kB+ FLASH



# Train and export a model



## 1. Train using standard Python ML libraries.

```
from keras import ...      A) keras neural  
                           network  
  
model = Sequential([  
    Dense(16, input_dim=n_features, activation='relu'),  
    Dense(8, activation='relu'),  
    Dense(1, activation='sigmoid'),  
])  
model.compile(...)  
model.fit(X_train, Y_train, epochs=1, batch_size=10)
```

```
from sklearn.neural_network import MLPClassifier  
model = MLPClassifier(hidden_layer_sizes=(100,50,25))  
  
model.fit(X_train, Y_train)  
  
B) scikit-learn neural network
```

## 2. Use `emlearn.convert()` and `.save()`

```
import emlearn  
  
cmodel = emlearn.convert(model, method='inline')  
  
cmodel.save(file=mynet_model.h', name='mynet')
```



```
#include <eml_net.h>  
  
static const float mynet_layer_0_biases[8] = { -0.015587f, -0.005395f, -0.010957f, 0.015883f ....  
static const float mynet_layer_0_weights[24] = { -0.256981f, 0.041887f, 0.063659f, 0.011013f, ....  
static const float mynet_layer_1_biases[4] = { 0.001242f, 0.010440f, -0.005309f, -0.006540f };  
static const float mynet_layer_1_weights[32] = { -0.577215f, -0.674633f, -0.376140f, 0.646900f, ....  
static float mynet_buf1[8];  
static float mynet_buf2[8];  
static const EmlNetLayer mynet_layers[2] = {  
    { 8, 3, mynet_layer_0_weights, mynet_layer_0_biases, EmlNetActivationRelu },  
    { 4, 8, mynet_layer_1_weights, mynet_layer_1_biases, EmlNetActivationSoftmax }  
};  
static EmlNet mynet = { 2, mynet_layers, mynet_buf1, mynet_buf2, 8 };  
  
int32_t  
mynet_predict(const float *features, int32_t n_features)  
{  
    return eml_net_predict(&mynet, features, n_features);  
}  
.....
```

Example of generated code

# Using the C code



## 3. #include and call predict()

```
// Include the generated model code  
#include "mynet_model.h"  
  
// index for the class we are  
detecting  
#define MYNET_VOICE 1  
  
// Buffers for input data  
#define N_FEATURES 6  
float features[N_FEATURES];  
  
#define DATA_LENGTH 128  
int16_t  
sensor_data[DATA_LENGTH];
```

setup

```
// Get data and pre-process it  
read_microphone(sensor_data, DATA_LENGTH);  
preprocess_data(sensor_data, features);  
  
// Run the model  
out = mynet_predict(features, N_FEATURES);  
  
// Do something with results  
if (out == MYNET_VOICE) {  
    set_display("voice detected");  
} else {  
    set_display("");  
}
```

loop



# Summary

## **1. Modern microcontrollers are very accessible**

Runs (Micro)Python!

ESP32 recommended

ViperIDE easy start

## **Conclusions**

## **2. MicroPython productive environment for sensor devices**

Python familiarity and ease-of-use

mip package manager

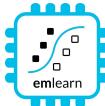
Good connectivity

## **3. Can implement advanced processing of sensor data**

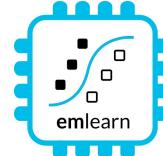
Accelerometer, audio, image, radar, ....

C modules a killer feature

emlearn-micropython: modules for DSP and Machine Learning



# More resources



emlearn-micropython

`emlearn_iir`  
`emlearn_trees`  
`emlearn_fft`  
`emlearn_cnn`  
`emlearn_neighbors`

Infinite Impulse Response filters  
Random Forest  
Fast Fourier Transform  
Convolutional Neural Networks  
K-nearest Neighbors

**Image classification+  
On-device learning**

**Ulab:** Numpy implementation for MicroPython

<https://github.com/v923z/micropython-ulab>

**OpenMV:** Computer/Machine Vision for MicroPython

<https://openmv.io/>

**PyCon Berlin 2024: Machine Learning on microcontrollers using MicroPython and emlearn**

<https://www.youtube.com/watch?v=S3GjLr0ZIE0>

**PyData ZA 2024: Sensor data processing on microcontrollers with MicroPython**

<https://za.pycon.org/talks/31-sensor-data-processing-on-microcontrollers-with-micropython/>

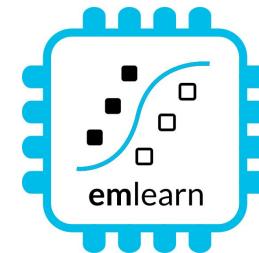
# Microcontrollers and Machine Learning

with MicroPython and emlearn

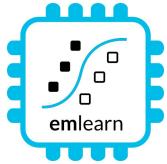
<https://github.com/emlearn/emlearn-micropython>



Jon Nordby [jononor@gmail.com](mailto:jononor@gmail.com)  
foss-north 2025



# Bonus



# Model optimization

[https://emlearn.readthedocs.io/en/  
latest/model\\_optimization.html](https://emlearn.readthedocs.io/en/latest/model_optimization.html)

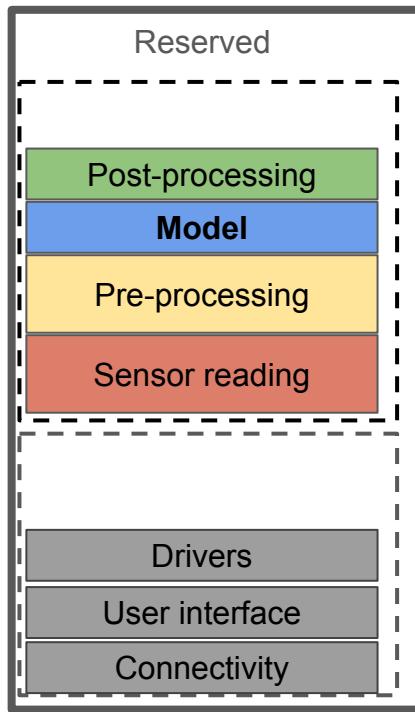
[https://emlearn.readthedocs.io/en/  
latest/tree\\_based\\_models.html](https://emlearn.readthedocs.io/en/latest/tree_based_models.html)

# Compute constraints

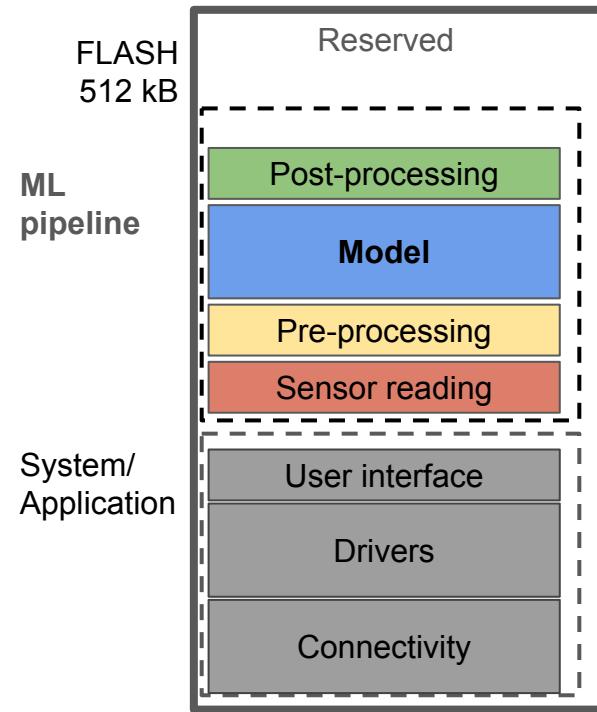
- Memory (RAM)
- Program space (FLASH)
- Inference time (CPU)
- Energy (battery)

CPU  
Cortex M4F

RAM  
64 kB



FLASH  
512 kB



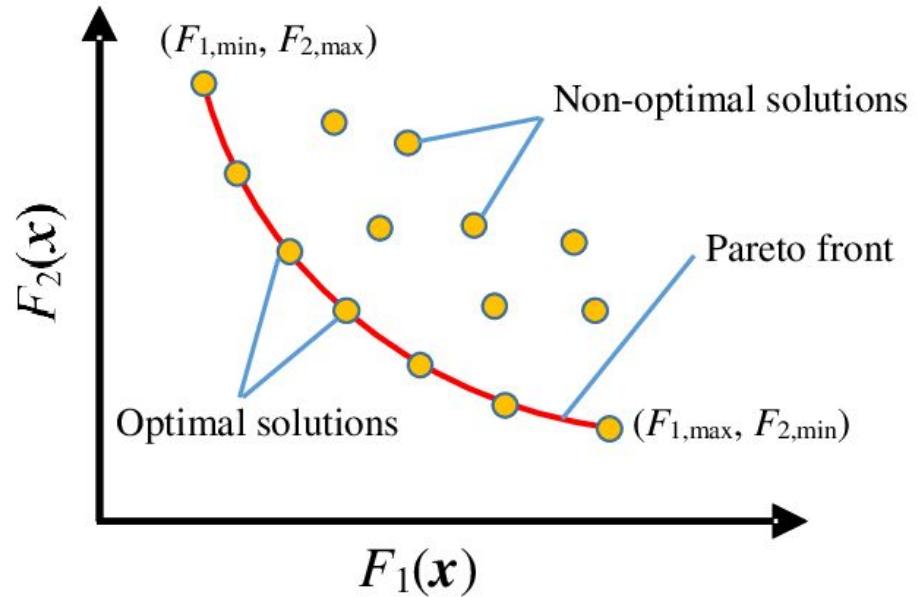
# Pareto front - performance/compute tradeoffs

Many possible combinations of **predictive performance vs computational cost**

No point considering the non-optimal solutions!

**Pareto front** is formed by the set of optimal solutions - that dominate the non-optimal ones

Describes the **possible tradeoffs between predictive performance and compute**



Utilities for finding Pareto front in `emlearn.evaluate.pareto`

<https://emlearn.readthedocs.io/en/latest/evaluate.html>

Example: [https://emlearn.readthedocs.io/en/latest/auto\\_examples/trees\\_hyperparameters.html](https://emlearn.readthedocs.io/en/latest/auto_examples/trees_hyperparameters.html)

# Tree-based ensembles - costs

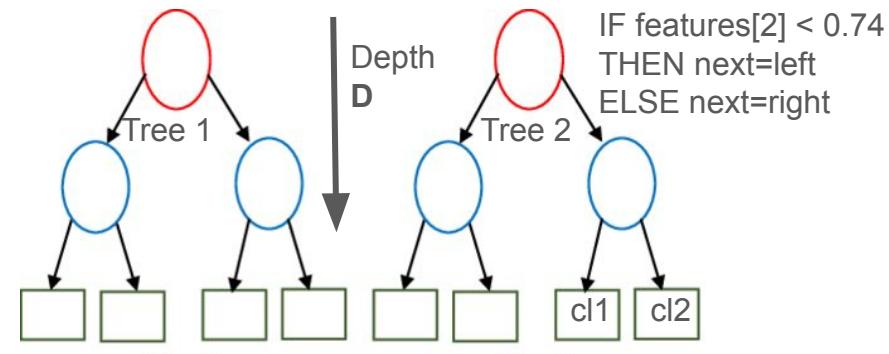
No. features  $F[0]$

No. trees  $T$

Tree Depth  $D_{\max}[1] / D_{\text{eff}}[2]$

No. nodes  $N \approx T * (2^{**} D)$

Input features, length  $F$   
[ 0.23, 0.56, 0.11, 0.55 ]



## Computational costs

Memory (RAM)  $O(F)$

Program (FLASH)  $O(N)$  [3]

Exc. time (CPU)  $\sum(D_{\text{eff}}(t), t \rightarrow T)$

## Utilities for estimating costs

```
from emlearn.evaluate.trees import  
model_size_bytes,  
compute_cost_estimate
```

[0] May also enable deeper trees

[1] Depth might differ across trees

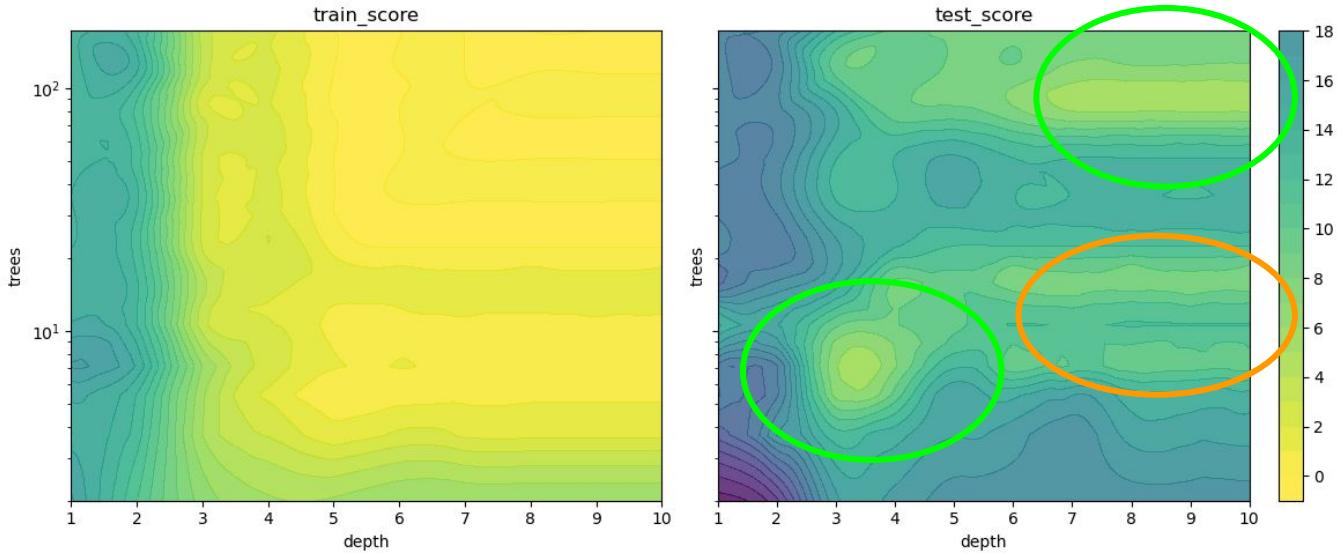
[2] Execution path is data dependent! Can estimate average using a dataset, or worst-case from model

[3] Assuming leaf and decision nodes same size

# Tree-based ensemble - predictive performance

No. features	$F$	increases capacity
No. trees	$T$	increases capacity, <b>decreases overfitting</b>
Tree Depth	$D_{\max} / D_{\text{eff}}$	increases capacity, <b>increases overfitting</b>
No. nodes	$N \approx T * D$	

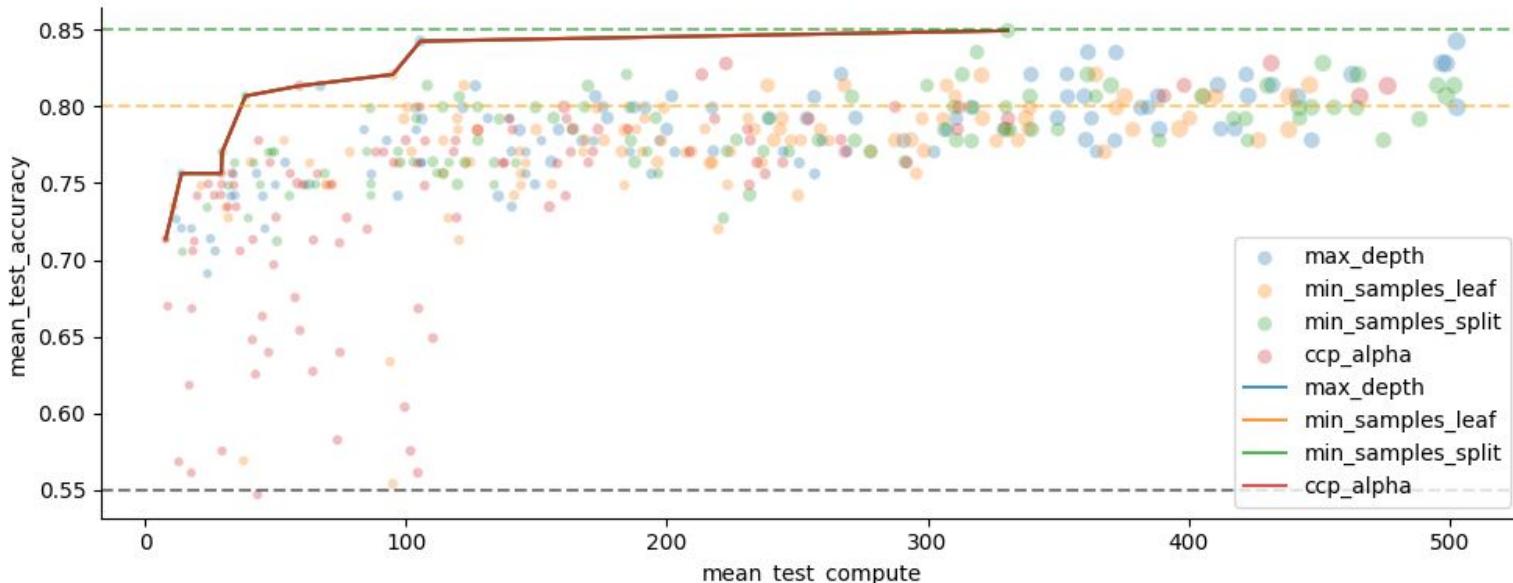
For **generalized performance** (on unseen data), **need to balance overfitting**.



# Hyperparameter tuning with scikit-learn

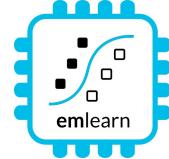
Vary **n\_estimators** (trees)  
and *one* of the **depth limiting** hyperparameters

from sklearn.model\_selection import **GridSearchCV**, **RandomizedSearchCV**



# Bonus

# TinyML for MicroPython - comparisons



Project	Deployment	Models	Size	Compute time
emlearn -micropython	Easy. Native mod .mpy	DT, RF, KNN, CNN	Good	Good
everywhereml	Easy. Pure Python .py	DT, RF, SVM, KNN,	High with large models	Poor
m2cgen	Easy. Pure Python .py	DT, RF, SVM, KNN, MLP	High with large models	Poor
OpenMV.tf	Hard. Custom Fork	CNN	High initial size	Good
ulab	Hard. User C module	(build-your-own) Using ndarray primitives	High initial size	Unknown (assume good)

# Microcontroller - tiny programmable chip

Compute power: 1 / 1000x of a smartphone

- RAM: 0.10 - 1 000 kB
- Program space: 1.0 - 10 000 kB
- Compute 10 - 1 000 DMIPS
- Price: 0.10 - 10 USD
- Energy: 1 - 1 000 milliWatt

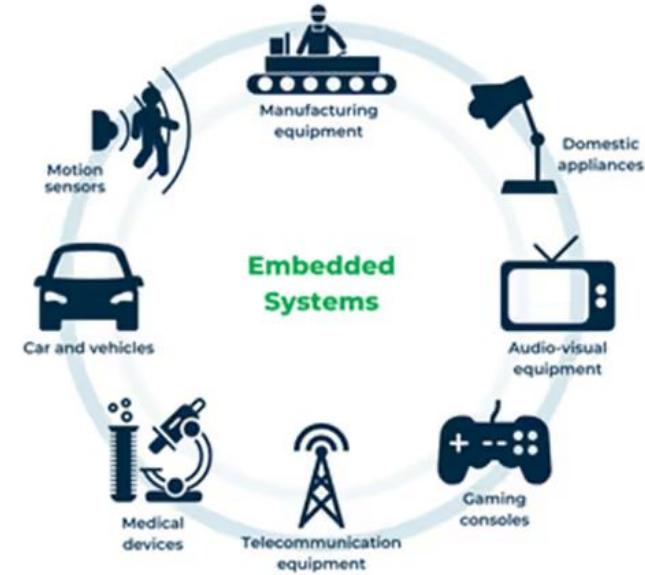
Over 20 *billions* shipped per year!

Increasingly accessible for hobbyists

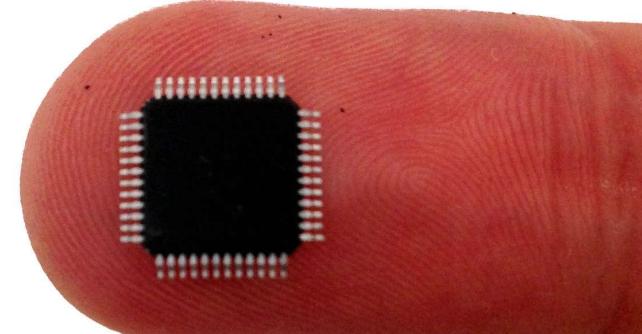
2010: Arduino Uno

2014: MicroPython

2019: MicroPython 1.10 - ESP32 PSRAM



Efficiency is key !  
Memory, compute, power



# Inline Assembly

MicroPython can expose Assembler opcodes as Python statements.

Allows to write a function in Assembler *inline in the Python program*

Can compile and execute on device

Supported on ARM Cortex M chips  
*Not supported (yet) on ESP32*

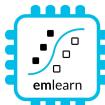
**For the most hardcore hackers!**

Official Documentation:

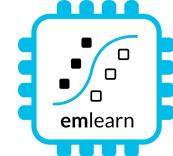
[https://docs.micropython.org/en/latest/reference/asm\\_thumb2\\_index.html](https://docs.micropython.org/en/latest/reference/asm_thumb2_index.html)

```
@micropython.asm_thumb
def fir(r0, r1, r2):
    mov(r3, r8)                      # For Pico: can't push({r8}). r0-r7 only.
    push({r3})
    ldr(r7, [r0, 0])                 # Array length
    mov(r6, r7)
    mov(r3, r0)
    add(r3, 12)                      # r3 points to ring buffer start
    sub(r7, 1)
    add(r7, r7, r7)
    add(r7, r7, r7)
    add(r5, r7, r3)
    ldr(r4, [r0, 8])
    cmp(r4, 0)
    bne(INITIALISED)
    mov(r4, r3)
    label(INITIALISED)
    str(r2, [r4, 0])
    add(r4, 4)
    cmp(r4, r5)
    ble(BUFOK)
    mov(r4, r3)
    label(BUFOK)
    str(r4, [r0, 8])                  # Save the insertion point for next call
    # *** Filter ***
    ldr(r0, [r0, 4])                  # Bits to shift
```

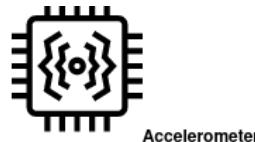
Example: FIR filter implementation (cut out)  
<https://github.com/peterhinch/micropython-filters/blob/master/fir.py>



# Sensor node systems



Including on-sensor data processing  
with Digital Signal Processing (DSP) and Machine Learning (ML)



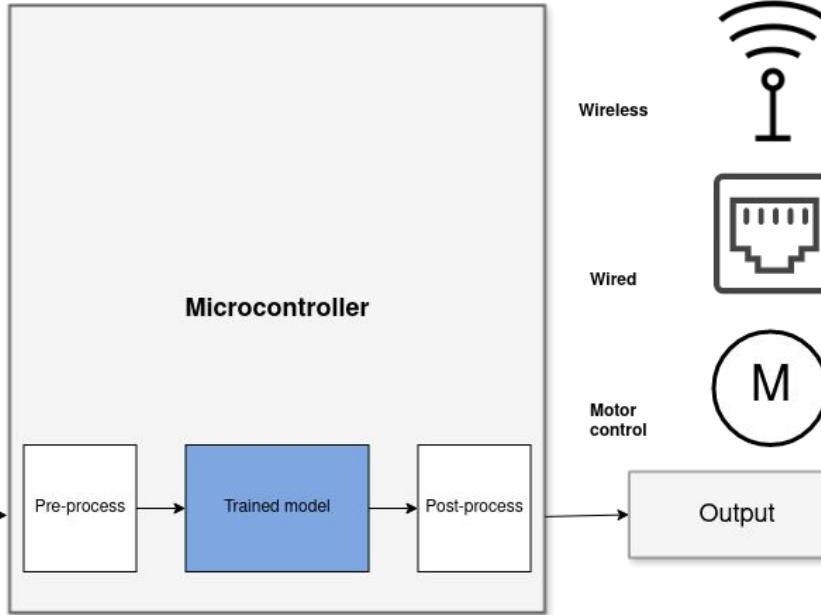
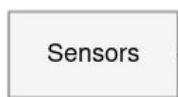
Accelerometer



Microphone



Camera



**“Internet of Things”**



**“Wearable devices”**



**Robotics,  
Industrial automation**

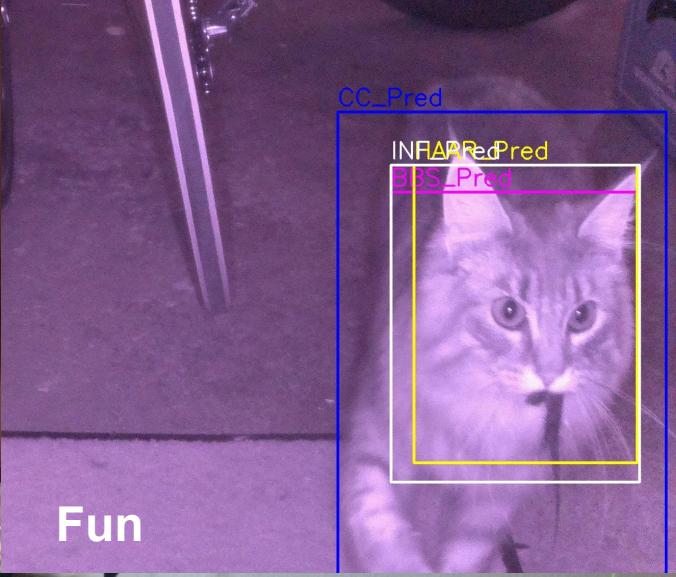
Hey Google...



Consumer Tech



Industrial



Fun



# Cat Detector - Data Acquisition

Classify *low-res* images, at 1 FPS+

Using mp\_camera by cnadler86  
[https://github.com/  
cnadler86/micropython-camera-API](https://github.com/cnadler86/micropython-camera-API)

96x96 px grayscale.  
Takes ~100 ms - OK



# Cat Detector - Classification with CNN

## emlearn\_tinyai\_x\_cnn

Go-to solution for image classification:  
Convolutional Neural Network (CNN)

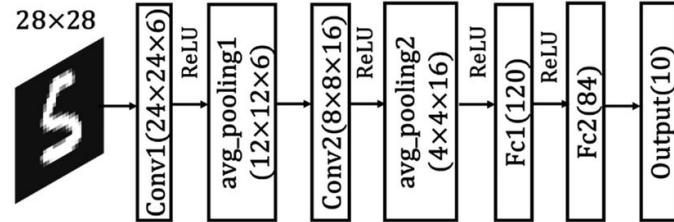
32x32 px input. 3 layers. Under 100 ms - OK

Preprocessing (untested)

Downscaling 96 px -> 32 px

Image brightness normalization

See also OpenMV - MicroPython for Computer Vision- <https://docs.openmv.io/index.html>



[https://github.com/emlearn/emlearn-micropython/tree/master/examples/mnist\\_cnn](https://github.com/emlearn/emlearn-micropython/tree/master/examples/mnist_cnn)

```
import tinyai_x_cnn # from emlearn-micropython

with open('cat_classifier.tmdl', 'rb') as f:
    model_data = array.array('B', f.read())
    model = tinyai_x_cnn.new(model_data)

while True:

    raw = read_camera()
    img = preprocess(raw)
    classification = model.predict(img)

    if classification == MY_CAT:
        open_door()

    machine.lightsleep(500)
```

